


A decorative graphic on the left side of the slide consisting of two overlapping parallelograms. The front one is blue and the back one is light green. They are positioned diagonally, with the blue one partially covering the green one.

Traffic Sign Classification CNN

By Alonso Bravo

Why Traffic Sign Classification?

- Autonomous Vehicles
- Vehicle Driving Assistance Features





Dataset Exploration

German Traffic Sign Recognition Benchmark (GTSRB)

- 30 Classes
- Training Set: 2607 Traffic Sign Images
- Testing Set: 1200 Traffic Sign Images
- Class Inconsistencies

ClassId, Name

0, Speed limit (5km/h)

1, Speed limit (15km/h)

2, Speed limit (30km/h)


3, Speed limit (40km/h)

4, Speed limit (50km/h)

5, Speed limit (60km/h)

6, Speed limit (70km/h)

7, speed limit (80km/h)





9 Random Images from Dataset



Image Pre-Processing

- Images Resized to 32x32 x3
- RGB Format
- Normalized Pixel Value: $\text{Pixel}/255$. (0 - 1)
 - Helps convergence in training



Image Diversity

- 10deg Random Rotation
 - 15% Random Zoom
 - 10% Horizontal Shift
 - 10% Vertical Shift
 - 15% Random Shearing
-
- Disabled flipping to ensure all traffic signs are read in the upright position
 - Gaps in image are filled with values from the nearest pixel



Actual=3 || Pred=3



Actual=3 || Pred=3



Actual=3 || Pred=3



Actual=3 || Pred=3



Actual=3 || Pred=4



Actual=3 || Pred=6



Actual=3 || Pred=3



Actual=3 || Pred=3



Actual=3 || Pred=3



Actual=3 || Pred=5



Actual=3 || Pred=6



Actual=3 || Pred=3



Actual=3 || Pred=3



Actual=3 || Pred=3



Actual=3 || Pred=3




Actual=3 || Pred=3



Actual=27 || Pred=27



Actual=27 || Pred=27



Actual=27 || Pred=27



Actual=27 || Pred=27



Actual=27 || Pred=27



Actual=27 || Pred=27



Actual=28 || Pred=28



Actual=28 || Pred=28



Actual=28 || Pred=28



Actual=28 || Pred=28



Actual=28 || Pred=28



Actual=28 || Pred=28



Actual=28 || Pred=28



Actual=28 || Pred=28



Actual=28 || Pred=28



Actual=28 || Pred=28

Pre-Processed Images



Model Architecture

- Sequential Convolutional Neural Network
 - Tensorflow / Keras
 - Sequential API
-
- Total Parameters: 5,260,604
 - Trainable Parameters: 1,753,086
 - Non-Trainable Parameters: 1,344

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 30, 30, 16)	448
conv2d_5 (Conv2D)	(None, 28, 28, 32)	4,640
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 32)	0
batch_normalization_3 (BatchNormalization)	(None, 14, 14, 32)	128
conv2d_6 (Conv2D)	(None, 12, 12, 64)	18,496
conv2d_7 (Conv2D)	(None, 10, 10, 128)	73,856
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 128)	0
batch_normalization_4 (BatchNormalization)	(None, 5, 5, 128)	512
flatten_1 (Flatten)	(None, 3200)	0
dense_2 (Dense)	(None, 512)	1,638,912
batch_normalization_5 (BatchNormalization)	(None, 512)	2,048
dropout_1 (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 30)	15,390

Layers and Parameters



Hyper Parameters

Learning Rate: 0.001

Optimizer: ADAM

Epochs: 30

Batch Size: 32

Layers: 13


- 4 Conv2D
- 2 Max Pool2D
- 3 Batch Normalization
- 1 Flatten
- 2 Dense
- 1 Dropout



Results

Test Data Accuracy: ~81%

	Precision	Recall	F1-Score
Macro	77%	80%	0.74
Weighted	88%	81%	0.81



`selected_classes = ['29', '1', '9', '25', '4', '5', '12', '7', '8', '13']`



Test Data Predictions starting at index: 675



How to Improve the Model?

- Larger and Higher Quality Dataset
- More Consistent Dataset
- More Robust Image Preprocessing

- Alternative:
 - Use a Pretrained Model (ResNet)



[https://drive.google.com/file/d/1BZYhN4VqWjrFCABZzW-Hgy6xpT19xw7N/view?usp=drive link](https://drive.google.com/file/d/1BZYhN4VqWjrFCABZzW-Hgy6xpT19xw7N/view?usp=drive_link)